

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
```

Matplotlib is building the font cache; this may take a moment.

```
In [2]: df = pd.read_csv('CarPrice_Assignment.csv')
df.head()
```

Out[2]:

	car_ID	symboling	CarName	fuelytype	aspiration	doornumber	carbody	drivewheel	enginel
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	

5 rows × 26 columns

```
In [3]: X = df[['enginesize', 'horsepower', 'citympg', 'highwaympg']]
y = df['price']
```

```
In [4]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s
```

```
In [5]: scaler= StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [6]: model = LinearRegression()
model.fit(X_train_scaled, y_train)
```

Out[6]:

▼ LinearRegression
LinearRegression()

```
In [7]: y_pred = model.predict(X_test_scaled)
```

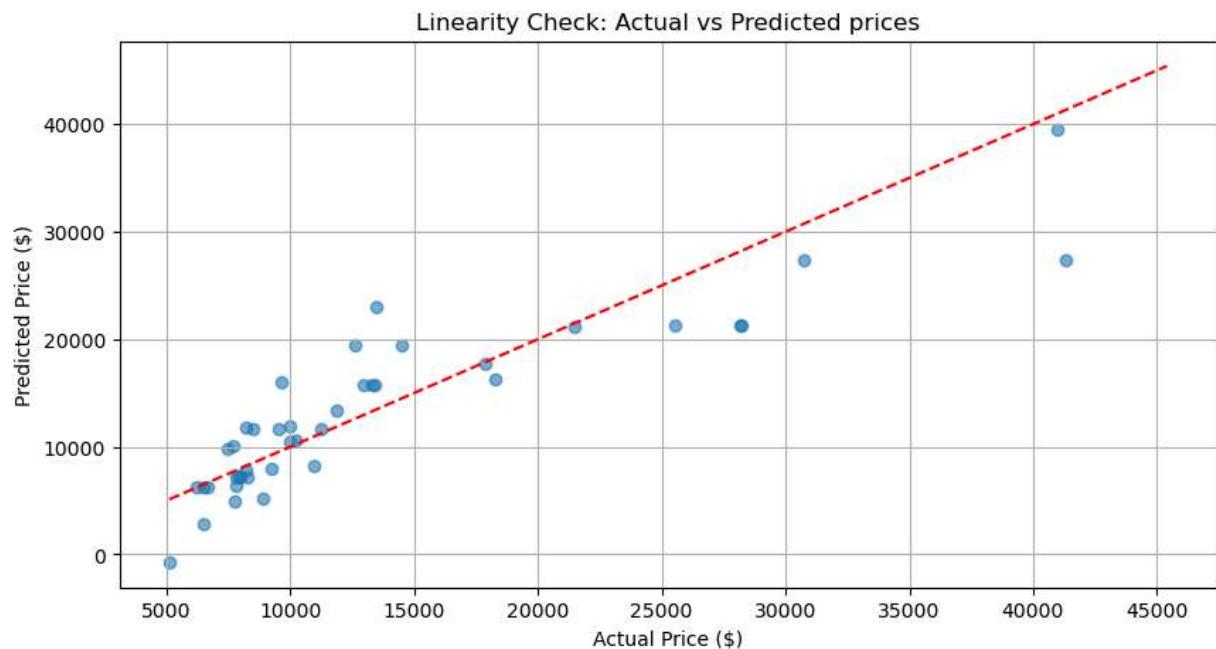
```
In [8]: print('Name: Raghul.S')
print('Reg. No: 212225040325')
print("MODEL COEFFICIENTS:")
for feature, coef in zip(X.columns, model.coef_):
    print(f'{feature}: {coef:.2f}')
print(f'{model.intercept_:.2f}')

print("\nMODEL PERFORMANCE:")
print(f'{mean_squared_error(y_test, y_pred):.2f}')
print(f'{np.sqrt(mean_squared_error(y_test, y_pred)):.2f}')
print(f'{r2_score(y_test, y_pred):.2f}')
print("*50")
```

Name: Raghul.S
Reg. No: 212225040325
MODEL COEFFICIENTS:
enginesize: 4523.40
horsepower: 1694.22
citympg: -392.57
highwaympg: -816.36
Intercept: 13223.41

MODEL PERFORMANCE:
MSE: 16471505.90
RMSE: 4058.51
R-squared: 0.79
=====

```
In [9]: plt.figure(figsize=(10, 5))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--')
plt.title("Linearity Check: Actual vs Predicted prices")
plt.xlabel("Actual Price ($)")
plt.ylabel("Predicted Price ($)")
plt.grid(True)
plt.show()
```



```
In [10]: residuals = y_test - y_pred
dw_test = sm.stats.durbin_watson(residuals)
print(f"\nDurbin-Watson Statistic: {dw_test:.2f},")
```

Durbin-Watson Statistic: 2.28

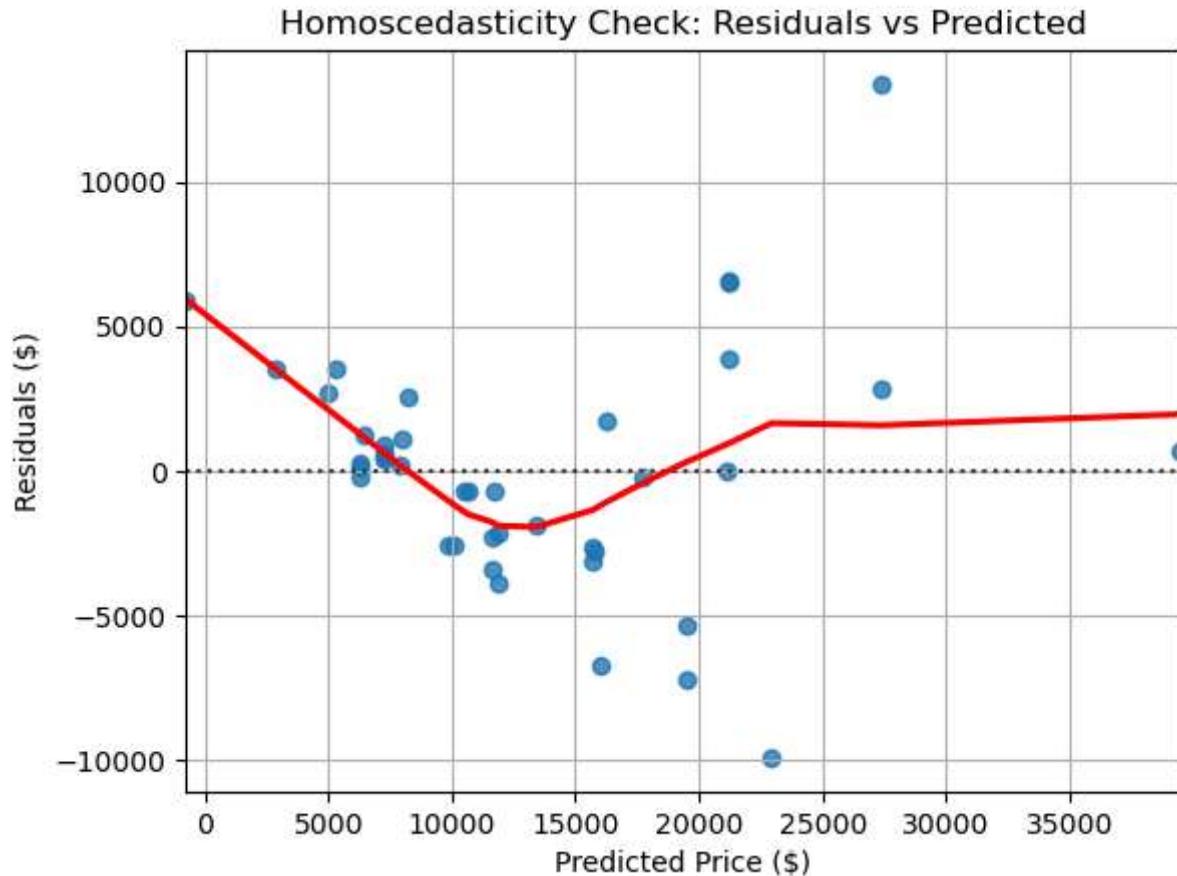
```
In [11]: residuals = y_test - y_pred
dw_test = sm.stats.durbin_watson(residuals)
print(f"\nDurbin-Watson Statistic: {dw_test:.2f}",
      "/n(Values close to 2 indicate no autocorrelation)")
```

Durbin-Watson Statistic: 2.28 /n(Values close to 2 indicate no autocorrelation)

```
In [12]: plt.figure(figsize=(10, 5))
```

```
Out[12]: <Figure size 1000x500 with 0 Axes>
<Figure size 1000x500 with 0 Axes>
```

```
In [13]: sns.residplot(x=y_pred, y=residuals, lowess=True, line_kws={'color': 'red'})
plt.title("Homoscedasticity Check: Residuals vs Predicted")
plt.xlabel("Predicted Price ($)")
plt.ylabel("Residuals ($)")
plt.grid(True)
plt.show()
```



```
In [15]: fig, (ax1, ax2) = plt.subplot(1, 2, figsize=(12, 5))
sns.histplot(residuals, kde=True, ax=ax1)
ax1.set_title("Residuals Distribution")
sm.qqplot(residuals, line='45', fit=True, ax=ax2)
ax2.set_title("Q-Q Plot")
```

NameError

Cell In[15], line 1

```
----> 1 fig, (ax1, ax2) = Plt.subplot(1, 2, figsize=(12, 5))
      2 sns.histplot(residuals, kde=True, ax=ax1)
      3 ax1.set_title("Residuals Distribution")
```

Traceback (most recent call last)

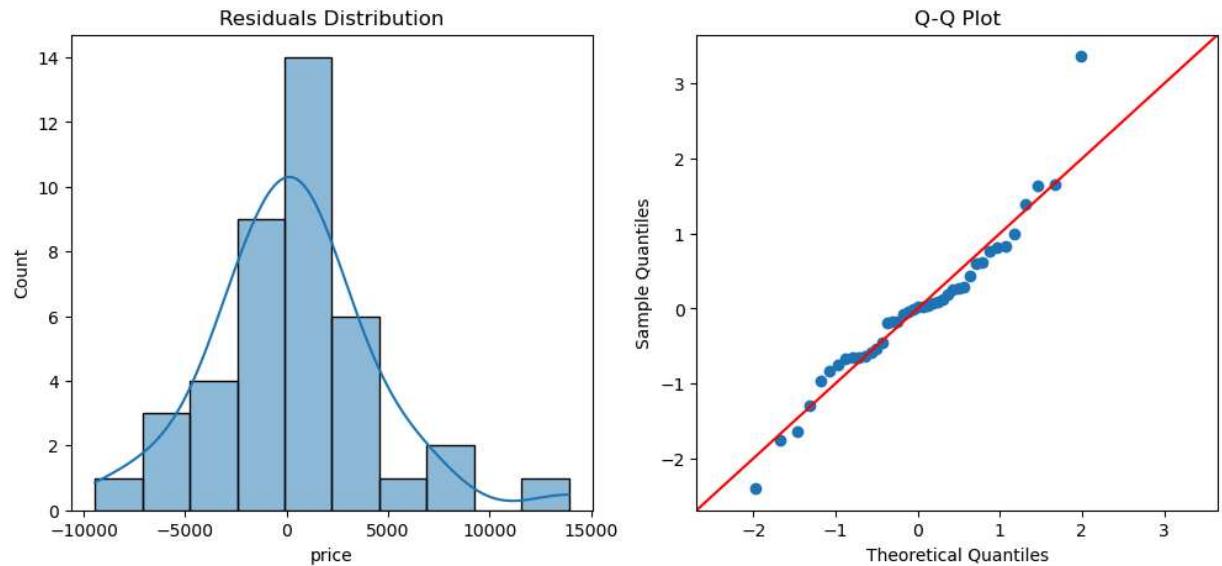
NameError: name 'Plt' is not defined

```
In [16]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))

sns.histplot(residuals, kde=True, ax=ax1)
ax1.set_title("Residuals Distribution")

sm.qqplot(residuals, line='45', fit=True, ax=ax2)
ax2.set_title("Q-Q Plot")
```

Out[16]: Text(0.5, 1.0, 'Q-Q Plot')



In []: